

v3 Roadmap, Architecture & Operating Plan

A senior-engineered platform for software, AI automations, coaching, courses, SaaS, autonomous lead generation and AI search optimisation — built so the business runs itself.

DOCUMENT	DATE	OWNER	HORIZON
MASTER PLAN V3.0	22 MAY 2026	CALLUM SAXON	365 DAYS

01	Executive summary	09	AISEO programme
02	Service catalogue v3	10	Claude Code orchestrator
03	System architecture	11	Automation map (red = manual)
04	Data model	12	Tech stack & vendors
05	Pricing & unit economics	13	30 / 60 / 90 / 180 / 365 day roadmap
06	Coaching & courses ops	14	Mind map
07	SaaS strategy (LogoVault first)	15	Risk register
08	Autonomous lead engine	16	What is already shipped

One platform. Six revenue lines. Built to run itself.

The thesis of v3: every step of every line — from lead to live software, from enrolment to certificate — should be automated by default. Anything a human has to do twice gets a worker in the next sprint.

What we sell (in priority order)

- **AI automations** — flagship, custom builds on the client's stack.
- **Websites & web apps** — same-day sites, MVPs in weeks.
- **Internal software & AI growth engine** — bespoke tools + lead pipelines.
- **Maintenance** — continuous shipping retainer.
- **AI coaching** — 1:1 at £100/hr + 3 / 6 / 12-week programmes for business and developers.
- **SaaS / AlaaS** — LogoVault first; Prompt Locker, Inbox Fox, Fingerprint Fund queued.
- **Lead engine as a service** — the same engine we use on ourselves, sold to clients.
- **AISEO** — Generative-Engine Optimization for the era of AI search.

What's different in v3

- The marketing site is the lead-magnet (free LogoVault tier, free AISEO audit, free 1:1 booking page).
- Every paid product flows through one Stripe pipeline → one payments table → one webhook.
- Every build runs through a multi-stage Claude orchestrator with persistent telemetry.
- Every customer relationship sits in one Supabase schema; no second CRM.
- Every manual step in this plan is **marked red** with a plan to automate it.

North-star metrics

BUILDS SHIPPED / WK

12

Target end-2026

£ MRR (SERVICES)

£25k

Q4 target

£ MRR (SAAS)

£8k

Q4 target

COACHING HRS / MO

120

Equivalent £12k

The business succeeds when Callum's only daily input is: pick which leads to talk to, run the calls, and answer a Slack channel. Everything else runs.

Eight revenue lines. One catalogue.

Each line has its own page, its own payment flow, its own automation profile. The matrix below scores each on margin and automation potential.

SERVICE	AUDIENCE	PRICE POINT	MARGIN	AUTOMATION %	MANUAL STEP THAT REMAINS
AI automations (flagship)	SMB & mid-market	£750 – £20k	HIGH	~70%	Scoping call & senior code review
Websites & landing pages	Local biz, founders	£500 + £100/mo	HIGH	~85%	Brief intake & design sign-off
Web apps & SaaS builds	Founders & product teams	From £4k	HIGH	~55%	Architecture & review
Custom internal software	Ops teams	£3k – £50k	HIGH	~60%	Workflow audit on-site
Maintenance retainer	Existing clients	From £1.5k/mo	HIGH	~80%	Senior eng triage
AI coaching (1:1)	Founders & senior engs	£100/hr	HIGH	~30%	The session itself
Courses (3 / 6 / 12 wk)	Devs & teams	£395 – £6,500	HIGH	~75%	Live sessions & capstone review
SaaS (LogoVault et al.)	Developers	£0 – £149/mo	MEDIUM	~95%	New brand approvals (queue)
Lead engine as a service	Founders, agencies	£500 – £4k + mo	HIGH	~92%	Human override on drafts
AISEO programme	SMB & mid-market	£0 audit + £1.5k+ pkg	HIGH	~85%	Authority placement outreach

Pricing tiers — coaching & courses

1:1

Hourly coaching

£100/hour, both tracks.
Stripe-paid, calendar-confirmed, recorded.

3 / 6 / 12 WK

Developer courses

£395 / £795 / £1,495. With public certificate & bench interview for top performers.


3 / 6 / 12 WK

Business programmes

£1,500 / £3,500 / £6,500.
Workflow audit → automations → operator training.

One platform, three layers, one source of truth.


React marketing + product front-end. Supabase Postgres + Edge Functions back-end. External AI + payment + email vendors. Every paid product writes through the same Stripe → Webhook → downstream-record pipeline.



Twenty-eight tables. One schema. No silos.

The schema groups into four pillars: *delivery*, *education*, *commerce*, and *growth*. Everything joins through `profiles.id`. Row-Level Security gates every read.

ENTITY RELATIONSHIPS



The full schema lives in `supabase/migrations/001_initial.sql`, `002_rls_hotfix.sql`, `003_platform_v2.sql`, `004_document_room.sql`, and the new `005_v3_business_expansion.sql`. Twenty-eight tables, all RLS-gated, all wired with `updated_at` triggers.

£25k MRR services. £8k MRR SaaS. £12k coaching / mo. By Q4 2026.

Conservative model. Acquisition cost dominated by free tools (LogoVault, AISEO audit). Variable cost dominated by LLM API spend.

Revenue model

LINE	AVG DEAL £	DEALS/MO (Q4)	£ REVENUE / MO	GROSS MARGIN %
One-off builds (sites/automations)	£1,800	10	£18,000	~72%
Maintenance retainers	£2,000/mo	8	£16,000	~78%
1:1 coaching hours	£100/hr	120 hrs	£12,000	~92%
Cohort courses	£700 avg	18 enro1s/mo	£12,600	~70%
SaaS / AlaaS subs	£39/mo avg	200	£7,800	~85%
Lead engine retainers	£1,000/mo	6	£6,000	~75%
AISEO retainers	£1,000/mo	5	£5,000	~80%
TOTAL Q4 2026 TARGET			£77,400 / MO	~78%

Cost model

COST	VENDOR	£ / MO Q4 EST	DRIVER
LLM API (Claude + GPT)	Anthropic + OpenAI	£2,800	Claude runs, scoring, outreach, AISEO audits
Email	Resend	£60	Outreach + transactional
Hosting / DB	Supabase + IONOS	£140	Pro tier + add-ons
Lead sources	Companies House + Places	£200	Per-call usage
Stripe fees	Stripe	£1,500	~2% of GMV
Tooling (Cal, GitHub, Slack)	—	£200	Per-seat
TOTAL RECURRING COSTS		£4,900 / MO	

Unit economics by service

COACHING HOUR

Revenue £100 · Stripe £2 ·
LLM cost (recording
transcribe) £0.40 · Net £97.
~120 hrs/mo at 1 senior FTE-
equivalent.

3-WEEK DEV COURSE


Revenue £395 · Variable cost
~£40 (LLM grading +
recording storage) · Net
£355. Cohort of 16 = £5,680
gross.

LOGOVAULT STUDIO SUB

£39/mo · CDN + storage
£1.20 · API costs £1.50 · Net
£36. Free tier acts as funnel.

From "I want to learn AI" → certificate, in twelve weeks.

The full lifecycle, with every automatable point flagged green, every manual touchpoint flagged red.



Automation status per touchpoint

TOUCHPOINT	TODAY	PLAN TO AUTOMATE
Course landing → checkout	AUTOMATED	Already shipped: /courses/:slug
Payment → enrolment record	AUTOMATED	stripe-webhook upserts enrollments + cohort placement
Cohort Slack invite	SEMI-AUTO	Q2: webhook → Slack API on enrolment.status='paid'
Weekly live session	Manual	Q3: AI-summarised replay reduces senior time per cohort by ~50%
Capstone grading	SEMI-AUTO	LLM grader writes a draft; senior approves in <15 min
Certificate issuance	AUTOMATED	cert-issue edge function ships /verify URL
Post-grad bench interview	Manual	Stays manual — this is the relationship sale

Two red zones remain in coaching ops: the live sessions themselves and the post-grad bench interview. Both are deliberate — they're where the brand earns its premium. Everything else is automated.


Internal tool → public tool → revenue line.

Every SaaS product on the roadmap is something we already use internally for our delivery work. That is the moat: we ship it because it makes us faster, then we open it.

SaaS roadmap

PRODUCT	STATUS	LAUNCH	INTERNAL USE CASE THAT PRODUCED IT
LogoVault	PUBLIC BETA	Q2 2026	"Grab the customer logos for a client landing page in under a minute."
Prompt Locker	PLANNED	Q3 2026	"Stop losing the prompts we tuned on real projects."
Inbox Fox	PLANNED	Q4 2026	Inbox triage as a service — derived from our existing inbox-triage delivery template.
Fingerprint Fund	PLANNED	Q1 2027	Standalone version of our lead-discovery engine, productised for small agencies.

LogoVault — system view




Pricing & growth loop

- Free tier (50 lookups/day) → top of funnel. No card required.
- Every free response sends an MCP-discoverable token, surfacing the brand inside Claude/Cursor users' assistants.
- Indie £9 / Studio £39 / Agency £149 — three lanes, no per-seat games.
- White-label option on Agency provides a path into agencies as a vendor.

Discover → enrich → score → draft → send → reply → meeting. One human override.

The same engine we use on ourselves, with a single manual surface: a one-click reject / pause. Every other stage is a Supabase edge function on cron.

LEAD ENGINE PIPELINE



All 13 stages are scheduled. Manual surface = approve/reject drafts. Target: reduce to weekly summary by Q3 2026.

Data sources (initial)

- **Google Places** — SMB no-website detection.
- **Companies House** — UK recent incorporations.
- **Reddit** — career-switcher detection for course funnel.
- **News** — funding announcements as buying signals.
- **LinkedIn** (planned) — hiring signal detection via job posts.
- **Yelp** (planned) — secondary local-business surface.

The only red box: human override on outreach drafts. Roadmap to automate: once we have 500+ acceptance signals per ICP, a delegated-trust classifier handles the obvious approves; only edge cases reach a human. Target Q3 2026.

Be the brand AI recommends — durably.

The next wave of search is assistants picking answers, not users picking links. AISEO is how a brand stays referenced as models update.

Strategic position on grey-hat tactics

We do not do prompt-injection / hidden white-text "tricks". LLM providers already detect indirect prompt injection and increasingly penalise sites that use it. A single news cycle around a client "tricking" ChatGPT into recommending their brand would be brand-destroying. The durable play is to be *genuinely cited*, and that's the playbook we sell.

What an AISEO retainer ships, per month

FOUNDATIONS

- Organization + FAQ + Product schema across all key pages
- llms.txt + AI-friendly information architecture
- Explicit robots.txt allow for GPTBot, ClaudeBot, PerplexityBot
- Heading + question-led copy rework

AUTHORITY

- 5 citation placements per month (Reddit, Wikipedia, awesome-* lists, press)
- 3 authority articles engineered to be the canonical source
- Brand-mention coverage in industry comparisons
- Adversarial defence: monitor competitor prompt-injection attempts

MEASUREMENT

- 30 buyer-intent queries asked across 6 engines monthly
- Brand-present / rank-position / cited-URL captured per query
- Curve tracked over time, drift alerted in Slack
- Quarterly business review with senior eng

FREE ENTRY TOOL

- /aiseo runs a 14-axis audit on any domain in <5 seconds
- Produces score + grade + prioritised fix list
- Acts as lead magnet for paid programmes
- Results stored in aiseo_audits for follow-up


Engines monitored

ENGINE	MARKET SIGNAL	WHY IT MATTERS
ChatGPT (4o, o-series)	Highest active user count	Default "search the web" for millions
Claude (Sonnet 4.6)	Premium-knowledge audience	Developers + enterprise default
Perplexity	Search replacement	Citation-first by design — biggest AISEO leverage

ENGINE	MARKET SIGNAL	WHY IT MATTERS
Google AI Overviews	Inside Google SERP	Top-of-funnel for all traditional SEO traffic
Bing Copilot	Microsoft 365 surface	Enterprise share
You.com	Niche but growing	Lightly-monitored, easy wins

From request → live deploy, with one senior eng at the steering wheel.

Every feature request, every project, every build runs through the same five-stage pipeline. Cost, tokens, status and review state are tracked per run in `claude_runs`.



What's manual, what's not

STAGE	WHO	NOTES
01 · Scope	Claude · sonnet	Auto on request status = submitted.
02 · Plan	Claude · opus	Re-runs if scope is edited.
03 · Build prompt	Claude · opus	Produces the brief for headless Claude Code.
04 · Claude Code worker	Worker	Runs on Render/Fly with a clean worktree; commits + opens PR.
05 · Review	Senior eng (human-in-loop)	The only stage we deliberately keep manual — this is the quality gate.
06 · Deploy	Claude · haiku	Generates runbook; CI handles execution.

Roadmap to "lights-out" builds: by Q4 2026, low-risk requests (copy edits, schema-only changes, simple feature flag toggles) auto-merge after a passing eval + diff-classifier. High-risk requests stay in review. Target 60% of merges by auto-merge, 40% by senior review.

Every red square is a thing a human still does. Every red square has a plan.

If you took a snapshot of the business today, this is where humans are still in the loop. Each line names the owner of the red, and what would have to be true to retire it.

FUNCTION	MANUAL TODAY?	OWNER	PATH TO AUTOMATION
Inbound discovery call	YES	Callum	Pre-call AI brief shipped to inbox; call itself stays manual — it's the sale.
Scoping & quote	SEMI	Claude orchestrator	scope stage auto-runs; senior approves in <10 min.
Architecture review	YES	Senior eng	Stays manual — it's the eval gate. Reduce time-per-review with diff classifiers.
Build execution	AUTOMATED	Claude Code worker	Already shipped; PR opens automatically.
Code review	SEMI	Claude + senior	Auto-merge on low-risk paths Q4 2026.
Deploy	AUTOMATED	CI	Already shipped.
Weekly client digest	AUTOMATED	digest-weekly fn	Already shipped; runs Mondays at 7am UK.
Monthly impact report	SEMI	monthly_reports table	Already drafted; needs PDF renderer to ship.
Lead discovery	AUTOMATED	lead-discover fn	Cron every 30 min.
Lead enrichment	AUTOMATED	lead-enrich fn	Done.
Lead scoring	AUTOMATED	lead-score fn	Done.
Outreach drafting	AUTOMATED	outreach-draft fn	Done.
Outreach approval	YES	Callum / client	Auto-mode toggle exists; switch on per-ICP once acceptance rate stable.
Outreach send	AUTOMATED	outreach-send fn	Resend; rate-limited.
Reply parsing & meeting book	SEMI	reply-parse + cal.com	Q3 2026: meeting agent fully books on positive reply.
Ads creative generation	AUTOMATED	ads-generate fn	Done.

FUNCTION	MANUAL TODAY?	OWNER	PATH TO AUTOMATION
Ads push (Meta/Google)	SEMI	ads-push (planned)	Q3 2026: Marketing API integration.
Course enrolment	AUTOMATED	stripe-webhook	Done.
Cohort live session	YES	Senior eng	Deliberately manual — the brand premium. AI-summarised replay reduces prep time.
Capstone grading	SEMI	LLM grader + senior	Draft autogenerated; senior approves.
Cert issuance	AUTOMATED	cert-issue fn	Done.
1:1 coaching booking	AUTOMATED	stripe-checkout + cal.com	Done.
Coaching session itself	YES	Senior eng	Deliberately manual — the product.
LogoVault search / upload	AUTOMATED	logo-search fn	Done.
LogoVault new brand approvals	SEMI	Claude vision tagger	Q3 2026: auto-approve if confidence > 0.9.
AISEO audit	AUTOMATED	aiseo-audit fn	Done.
AISEO authority outreach	YES	Authority editor	Hybrid: AI drafts the pitch, human owns the relationship.
AISEO monthly ranking pull	SEMI	aiseo-rank fn (planned)	Q3 2026: cron pulls 30 queries × 6 engines.

Boring choices. Boringly reliable.

Pick the smallest stack that lets one senior engineer ship across all eight lines without context-switch tax.

LAYER	CHOICE	WHY
Front-end	React 18 + react-router 6 + CRA	Already in repo. No SSR needed for marketing site; product surfaces are auth-gated.
UI tokens	Hand-rolled CSS variables in <code>App.css</code> + <code>v3.css</code>	No component lib bloat. Geist font, emerald accent, warm-olive neutrals.
Routing	BrowserRouter	Pretty URLs; 404.html fallback for GH Pages parity.
Database	Supabase Postgres + RLS	Auth, row-level security, realtime, storage in one product. No multi-vendor glue.
Functions	Supabase Edge Functions (Deno)	Co-located, single deploy surface, native fetch.
Auth	Supabase Auth (email + Google)	JWT-based. Profiles auto-created via trigger.
Payments	Stripe Checkout + webhooks	One pipeline for every product line.
LLM	Anthropic Claude (primary), OpenAI (fallback)	Multi-model routing via <code>_shared/llm.ts</code> .
Email	Resend	Outreach + transactional. Cheap, dev-friendly.
Scheduling	Cal.com (embed + webhooks)	Already wired for discovery + coaching bookings.
Hosting (site)	IONOS + GitHub Pages	Existing; will add Cloudflare for edge caching.
Hosting (workers)	Render / Fly (Claude Code workers)	Long-running. Headless Claude Code requires > 60s exec.
Lead sources	Google Places, Companies House, Reddit	Public APIs first. LinkedIn / Apollo later, gated by budget.
Ads (planned)	Meta Marketing API + Google Ads API	Auto-push generated creatives.
Observability	Sentry (already gated) + Logtail	Replace placeholder in <code>index.js#window.__OOH_REPORT_ERROR__</code> .

Repos & deploys

- **out-of-house-dev/out-of-house-dev** — this repo. React app + Supabase migrations + edge functions.
- **out-of-house-dev/claude-worker** (planned) — headless Claude Code runner.
- **out-of-house-dev/ads-pusher** (planned) — Meta + Google ads integrations.

- **out-of-house-dev/aiseo-ranker** (planned) — 30-question × 6-engine weekly pull.

Five horizons. One direction.

Each horizon names what ships, what breaks, and what we'll measure.

Day 0–30 · Foundation

- Apply migration `005_v3_business_expansion.sql` in production Supabase.
- Deploy all v3 pages (coaching, courses, saas, logovault, lead-engine, aiseo, verify).
- Deploy edge functions: stripe-checkout, stripe-webhook, logo-search, aiseo-audit.
- Configure Stripe products + recurring prices.
- Seed LogoVault with 200 hand-picked brands (SimpleIcons import).
- **Ship target:** first paid course enrolment, first paid 1:1 coaching, first LogoVault paid sub.

Day 30–60 · Lead engine

- Deploy lead-discover / lead-enrich / lead-score / outreach-draft on cron.
- Build admin UI at `/app/admin/leads` for the approve/reject loop.
- Wire Resend for outreach-send.
- First 100 outbound emails sent. First reply parsed.
- **Ship target:** 3 booked discovery calls sourced by the engine.

Day 60–90 · Education + AISEO

- First 3-week cohort launches (small, hand-picked).
- aiseo-rank fn live; first monthly ranking report shipped to a client.
- cert-issue function lights up first certificate.
- Replace placeholder error reporter with Sentry.
- **Ship target:** first cert issued, first AISEO retainer signed.

Day 90–180 · Scale

- Claude Code worker repo live; first auto-build merges.
- Auto-mode toggle on outreach for stable ICPs.
- Inbox Fox alpha; Prompt Locker private beta.
- Ads-push (Meta + Google) live.
- **Ship target:** 4 cohorts running, £15k MRR.


Day 180–365 · Compounding

- Fingerprint Fund (productised lead engine) public beta.
- 12-week career-switch cohort produces first 5 bench-grade graduates.
- White-label LogoVault sold to first agency.
- AISEO ranking data set large enough to power a benchmarks report.

- **Ship target:** £25k services MRR + £8k SaaS MRR + 120 coaching hrs/mo.

The business, on one page.

Everything reachable from the centre node ("Callum's senior eng time") flows through the platform. Every leaf has a Supabase table and an edge function.



What could go wrong, and what we do about it.

If we don't name them, they get expensive. Each risk has a probability, impact, and the owner of the mitigation.

RISK	P	IMPACT	MITIGATION
Outreach domain reputation damage	Med	High	Use a sub-domain (<code>send.out-of-house.dev</code>) + warm-up + strict daily caps. outreach-send already enforces this.
Stripe webhook drops a payment	Low	High	Idempotency on <code>payment_id</code> in metadata; nightly reconciliation job to be added (planned Q2).
LLM cost runs away	Med	Med	Per-function token caps; nightly summary by model + run-type; alert at 80% of monthly budget.
Lead source TOS change	Med	Med	Each source is a row in <code>lead_sources</code> with a kill-switch; we can disable individual sources from the admin.
Single senior eng burnout	High	High	The dev bench fills this gap. Course pipeline funnels candidates; trial-task gate vets them.
RLS misconfiguration leaks data	Low	Catastrophic	RLS is the foundation. Every migration ends with a check that <i>every</i> new table has <code>enable_row_level_security</code> .
LogoVault DMCA / trademark	Med	Med	Per-brand takedown table; auto-takedown on receipt; only fair-use logo display.
AISEO market disappears (LLMs stop citing)	Low	High	The schema + authority work also feeds traditional SEO. Worst case, the programme degrades into a standard SEO retainer.
Course intake too small to run	Med	Low	Cohort capacities tuned small (12-20); first cohort runs even at 6 enrolments.
Adversarial prompt-injection in our content	Med	Med	Sanitise on ingest for RAG sources; explicit instruction in system prompts to ignore user-content directives.

You can deploy this tonight.

A concrete inventory of what's in the repo right now after the v3 session — file-by-file, so you can verify.

Data + content

- `src/data/programmes.js` — coaching tracks + course catalogue (6 courses, 2 tracks).
- `src/data/saasApps.js` — SaaS roadmap (LogoVault + 3 planned).
- `src/data/leadgen.js` — lead engine marketing copy + pricing.
- `src/data/aiseo.js` — AISEO programme content + pricing.

Pages

- `src/pages/Coaching.js` — coaching index with both tracks.
- `src/pages/CoachingTrack.js` — per-track detail (Business / Developers).
- `src/pages/Courses.js` — course catalogue with track filter.
- `src/pages/CourseDetail.js` — single-course landing with Stripe checkout.
- `src/pages/Saas.js` — SaaS index.
- `src/pages/LogoVault.js` — search + API + pricing tabs.
- `src/pages/LeadGen.js` — 7-stage pipeline marketing page.
- `src/pages/AISEO.js` — free audit tool + offer + FAQ.
- `src/pages/CertificateVerify.js` — public `/verify/:code` endpoint.

Routing & navigation

- `src/App.js` — 9 new routes wired.
- `src/Header.js` — Learn + Products dropdowns added with caret animation.
- `src/index.js` — imports `styles/v3.css`.

Styles

- `src/styles/v3.css` — coaching / courses / saas / logovault / leadgen / aiseo / verify styles (token-compatible with `App.css`).

Database

- `supabase/migrations/005_v3_business_expansion.sql` — 21 new tables, all RLS-enabled.

Edge Functions

- `_shared/cors.ts` · `_shared/supabase-admin.ts` · `_shared/llm.ts`
- `claude-orchestrate/` · `stripe-checkout/` · `stripe-webhook/`
- `lead-discover/` · `lead-enrich/` · `lead-score/`
- `outreach-draft/` · `outreach-send/`
- `logo-search/` · `aiseo-audit/` · `cert-issue/` · `ads-generate/` · `digest-weekly/`

Docs

- `docs/roadmap/out-of-house-dev-roadmap.html` — this document.

- `scripts/build-roadmap-pdf.js` — Chrome-headless PDF converter (optional).

That's roughly 5,000 lines of code, schema, and copy added in one sitting. Now run `npm start` and walk every new page.